


Interpretable Machine Learning

Partial Dependence (PD) plot



Learning goals

- PD plots and relation to ICE plots
- Interpretation of PDP
- Extrapolation and Interactions in PDPs
- Centered ICE and PDP

PARTIAL DEPENDENCE (PD)

► Friedman (2001)

Definition: PD function is expectation of $\hat{f}(\mathbf{x}_S, \mathbf{x}_{-S})$ w.r.t. marginal distribution of features \mathbf{x}_{-S} :

$$f_{S,PD}(\mathbf{x}_S) = \mathbb{E}_{\mathbf{x}_{-S}} \left(\hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}) \right) = \int_{-\infty}^{\infty} \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}) d\mathbb{P}(\mathbf{x}_{-S})$$

Estimation: For a grid value \mathbf{x}_S^* , average ICE curves point-wise at \mathbf{x}_S^* over all observed $\mathbf{x}_{-S}^{(i)}$:

$$\hat{f}_{S,PD}(\mathbf{x}_S^*) = \frac{1}{n} \sum_{i=1}^n \hat{f}(\mathbf{x}_S^*, \mathbf{x}_{-S}^{(i)})$$

EXAMPLE: PARTIAL DEPENDENCE FOR LINEAR MODEL

Assume a linear regression model with two features:

$$\hat{f}(\mathbf{x}) = \hat{f}(\mathbf{x}_1, \mathbf{x}_2) = \hat{\beta}_1 \mathbf{x}_1 + \hat{\beta}_2 \mathbf{x}_2 + \hat{\beta}_0$$

PD function for feature of interest $S = \{1\}$ (with $-S = \{2\}$) is:

$$\begin{aligned} f_{1,PD}(\mathbf{x}_1) &= \mathbb{E}_{\mathbf{x}_2} \left(\hat{f}(\mathbf{x}_1, \mathbf{x}_2) \right) = \int_{-\infty}^{\infty} \left(\hat{\beta}_1 \mathbf{x}_1 + \hat{\beta}_2 \mathbf{x}_2 + \hat{\beta}_0 \right) d\mathbb{P}(\mathbf{x}_2) \\ &= \hat{\beta}_1 \mathbf{x}_1 + \hat{\beta}_2 \cdot \int_{-\infty}^{\infty} \mathbf{x}_2 d\mathbb{P}(\mathbf{x}_2) + \hat{\beta}_0 \\ &= \hat{\beta}_1 \mathbf{x}_1 + \underbrace{\hat{\beta}_2 \cdot \mathbb{E}_{\mathbf{x}_2}(\mathbf{x}_2)}_{:=const} + \hat{\beta}_0 \end{aligned}$$

\Rightarrow PD plot visualizes the function $f_{1,PD}(\mathbf{x}_1) = \hat{\beta}_1 \mathbf{x}_1 + const$ ($\hat{=}$ feature effect of \mathbf{x}_1).


PARTIAL DEPENDENCE

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	1	4	7	0.4
2	1	5	8	0.6
3	1	6	9	0.1

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	2	4	7	0.6
2	2	5	8	0.8
3	2	6	9	0.5

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	3	4	7	0.7
2	3	5	8	0.9
3	3	6	9	0.6

$\frac{1}{3} \sum_{i=1}^3 \hat{f}$
$\frac{1}{3} (0.4 + 0.6 + 0.1)$
$\frac{1}{3} (0.6 + 0.8 + 0.5)$
$\frac{1}{3} (0.7 + 0.9 + 0.6)$



Estimate PD function by **point-wise** average of ICE curves at grid value $\mathbf{x}_s^* = \mathbf{x}_1^* = 1$:

$$\hat{f}_{1,PD}(\mathbf{x}_1^*) = \frac{1}{n} \sum_{i=1}^n \hat{f}(\mathbf{x}_1^*, \mathbf{x}_{2,3}^{(i)})$$


PARTIAL DEPENDENCE

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	1	4	7	0.4
2	1	5	8	0.6
3	1	6	9	0.1

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	2	4	7	0.6
2	2	5	8	0.8
3	2	6	9	0.5

$\frac{1}{3} \sum_{i=1}^3 \hat{f}$
$\frac{1}{3} (0.4 + 0.6 + 0.1)$
$\frac{1}{3} (0.6 + 0.8 + 0.5)$
$\frac{1}{3} (0.7 + 0.9 + 0.6)$

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	3	4	7	0.7
2	3	5	8	0.9
3	3	6	9	0.6



Estimate PD function by **point-wise** average of ICE curves at grid value $\mathbf{x}_s^* = x_1^* = 2$:

$$\hat{f}_{1,PD}(x_1^*) = \frac{1}{n} \sum_{i=1}^n \hat{f}(x_1^*, \mathbf{x}_{2,3}^{(i)})$$


PARTIAL DEPENDENCE

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	1	4	7	0.4
2	1	5	8	0.6
3	1	6	9	0.1

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	2	4	7	0.6
2	2	5	8	0.8
3	2	6	9	0.5

i	\mathbf{x}_s	\mathbf{x}_2	\mathbf{x}_3	\hat{f}
1	3	4	7	0.7
2	3	5	8	0.9
3	3	6	9	0.6

$\frac{1}{3} \sum_{i=1}^3 \hat{f}$
$\frac{1}{3} (0.4 + 0.6 + 0.1)$
$\frac{1}{3} (0.6 + 0.8 + 0.5)$
$\frac{1}{3} (0.7 + 0.9 + 0.6)$



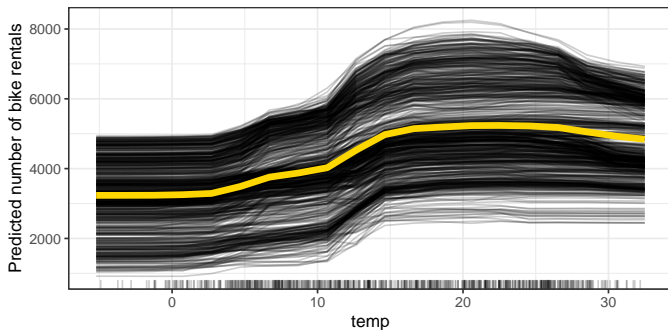
Estimate PD function by **point-wise** average of ICE curves at grid value $\mathbf{x}_s^* = \mathbf{x}_1^* = 3$:

$$\hat{f}_{1,PD}(x_1^*) = \frac{1}{n} \sum_{i=1}^n \hat{f}(x_1^*, \mathbf{x}_{2,3}^{(i)})$$

INTERPRETATION: PD AND ICE

If feature varies:


- **ICE:** How does **prediction of individual observation** change? \Rightarrow **local** interpretation
- **PD:** How does **average effect / expected prediction** change? \Rightarrow **global** interpretation



INTERPRETATION: PD AND ICE

If feature varies:

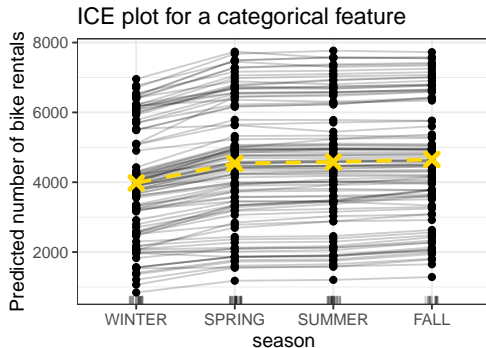

- **ICE:** How does **prediction of individual observation** change? \Rightarrow **local** interpretation
- **PD:** How does **average effect / expected prediction** change? \Rightarrow **global** interpretation



Insights from bike sharing data:

- Parallel ICE curves = homogeneous effect
- Warmer \Rightarrow more rented bikes
- Too hot \Rightarrow slightly less bikes


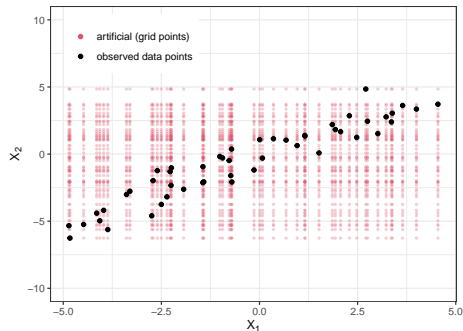
INTERPRETATION: CATEGORICAL FEATURES



- PDP with boxplots and ICE with parallel coordinates plots
- NB: Categories can be unordered, if so, rather compare pairwise

COMMENTS ON EXTRAPOLATION


Extrapolation can cause issues in regions with few observations or if features are correlated



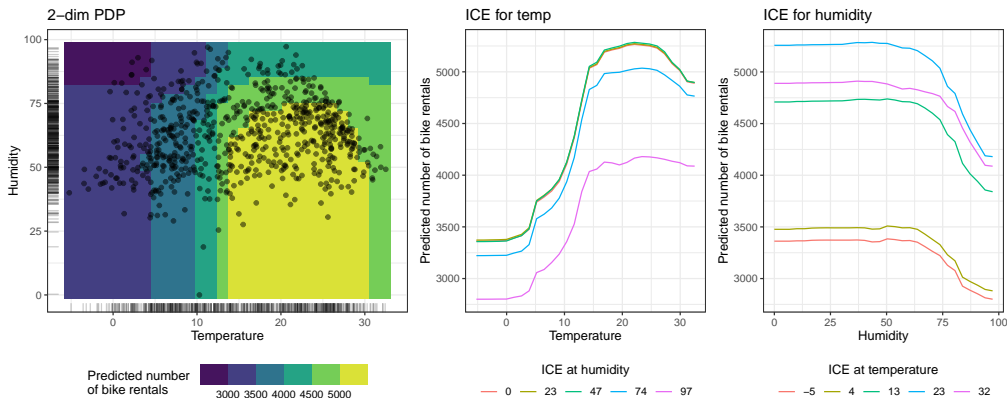
- **Example:** Features x_1 and x_2 are strongly correlated
- **Black points:** Observed points of the original data
- **Red:** Grid points used to calculate the ICE and PD curves (several unrealistic values)
 - ⇒ PD plot at $x_1 = 0$ averages predictions over the whole marginal distribution of feature x_2
 - ⇒ May be problematic if model behaves strange outside training distribution

COMMENTS ON INTERACTIONS

- PD plots: averaging of ICE curves might **obfuscate** heterogeneous effects and interactions
 - ⇒ Ideally plot ICE curves and PD plots together to uncover this fact
 - ⇒ Different shapes of ICE curves suggest interaction (but does not tell with which feature)



COMMENTS ON INTERACTIONS - 2D PARTIAL DEPENDENCE



- Humidity and temperature interact with each other at high values (see shape difference)
 \rightsquigarrow Shape of ICE curves at different horizontal and vertical slices varies (for high values)
- Low to medium humidity and high temperature \Rightarrow many rented bikes

CENTERED ICE PLOT (C-ICE)

Issue: Difficult to identify heterogenous ICE curves if curves have different intercepts (are stacked)

Solution: Center ICE curves at fixed reference value $x' \sim \mathbb{P}(\mathbf{x}_S)$, often $x' = \min(\mathbf{x}_S)$

⇒ Easier to identify heterogenous shapes with c-ICE curves

$$\begin{aligned}\hat{f}_{S,cICE}^{(i)}(\mathbf{x}_S) &= \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}^{(i)}) - \hat{f}(x', \mathbf{x}_{-S}^{(i)}) \\ &= \hat{f}_S^{(i)}(\mathbf{x}_S) - \hat{f}_S^{(i)}(x')\end{aligned}$$

⇒ Visualize $\hat{f}_{S,cICE}^{(i)}(\mathbf{x}_S^*)$ vs. grid point \mathbf{x}_S^*

CENTERED ICE PLOT (C-ICE)

Issue: Difficult to identify heterogenous ICE curves if curves have different intercepts (are stacked)





Solution: Center ICE curves at fixed reference value $x' \sim \mathbb{P}(\mathbf{x}_S)$, often $x' = \min(\mathbf{x}_S)$

⇒ Easier to identify heterogenous shapes with c-ICE curves

$$\begin{aligned}\hat{f}_{S, cICE}^{(i)}(\mathbf{x}_S) &= \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}^{(i)}) - \hat{f}(x', \mathbf{x}_{-S}^{(i)}) \\ &= \hat{f}_S^{(i)}(\mathbf{x}_S) - \hat{f}_S^{(i)}(x')\end{aligned}$$


⇒ Visualize $\hat{f}_{S, cICE}^{(i)}(\mathbf{x}_S^*)$ vs. grid point \mathbf{x}_S^*

Interpretation (yellow curve in c-ICE):
On average, the number of bike rentals at $\sim 97\%$ humidity decreased by 1000 bikes compared to a humidity of 0 %



CENTERED ICE PLOT (C-ICE)

For categorical features, c-ICE plots can be interpreted as in LMs due to reference value



Interpretation:

- The reference category is $x' = \text{SPRING}$
- Golden crosses: Average number of bike rentals if we jump from SPRING to any other season
 \Rightarrow Number of bike rentals drops by ~ 560 in WINTER and is slightly higher in SUMMER and FALL compared to SPRING