Interpretable Machine Learning

Inherently Interpretable Models - Motivation

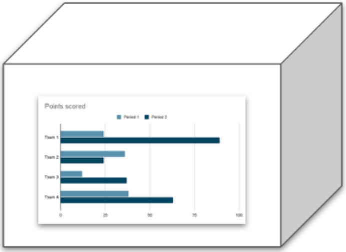

Learning goals

- Why should we use interpretable models?
- Advantages and disadvantages of interpretable models

MOTIVATION

- Achieving interpretability by using interpretable models is the most straightforward approach
- Classes of models deemed interpretable:
  - (Generalized) linear models (LM, GLM)
  - Generalized additive models (GAM)
  - Decision trees
  - Rule-based learning
  - Model-based boosting / component-wise boosting
  - ...

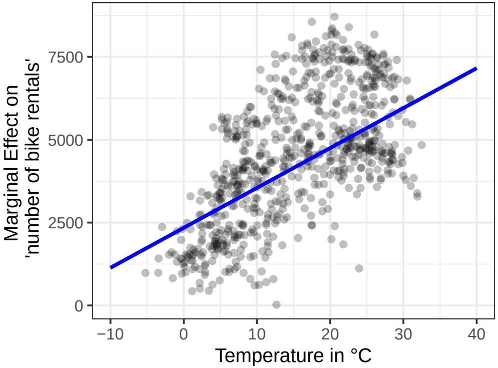

$\rightsquigarrow$ LM provides straightforward interpretation

ADVANTAGES

- Inherently interpretable - some additional model-agnostic interpretation methods not required
  $\rightsquigarrow$ Eliminates a source of error

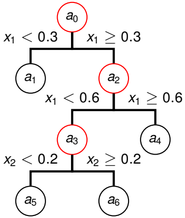

ADVANTAGES

- Inherently interpretable - some additional model-agnostic interpretation methods not required
  $\leadsto$ Eliminates a source of error

- Interpretable models often simple
  $\leadsto$ training time is fairly small

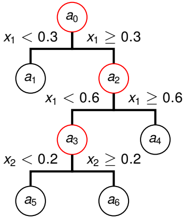

ADVANTAGES

- Inherently interpretable - some additional model-agnostic interpretation methods not required
  $\rightsquigarrow$ Eliminates a source of error
- Interpretable models often simple
  $\rightsquigarrow$ training time is fairly small
- Some interpretable models fulfill the monotonicity constraint
  $\rightsquigarrow$ Larger feature values always lead to higher (or smaller) outcomes (e.g., GLMs)

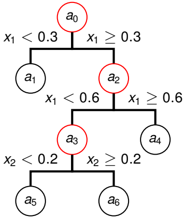

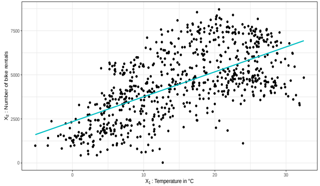

ADVANTAGES

- Inherently interpretable - some additional model-agnostic interpretation methods not required
  $\rightsquigarrow$ Eliminates a source of error

- Interpretable models often simple
  $\rightsquigarrow$ training time is fairly small

- Some interpretable models fulfill the monotonicity constraint
  $\rightsquigarrow$ Larger feature values always lead to higher (or smaller) outcomes (e.g., GLMs)

- Many people are familiar with traditional interpretable models
  $\rightsquigarrow$ Increases trust, facilitates communication of results

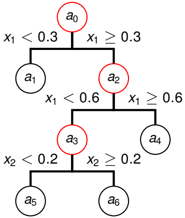

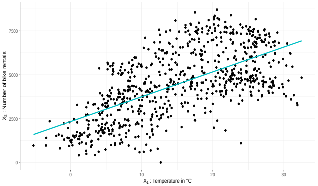

DISADVANTAGES

- Often certain assumptions about data / model structure required
  $\rightsquigarrow$ If assumptions are wrong, models may perform bad

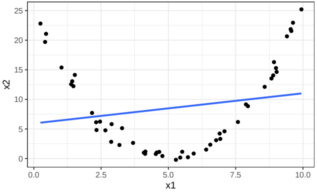

DISADVANTAGES

- Often certain assumptions about data / model structure required
  $\rightsquigarrow$ If assumptions are wrong, models may perform bad

- Interpretable models may also be difficult to interpret, e.g.:
  - Linear model with hundreds of features and interactions
  - Decision trees with huge tree depth

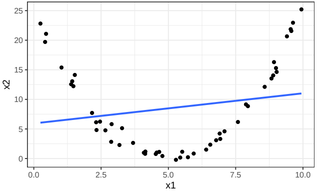

DISADVANTAGES

- Often certain assumptions about data / model structure required
  $\rightsquigarrow$ If assumptions are wrong, models may perform bad

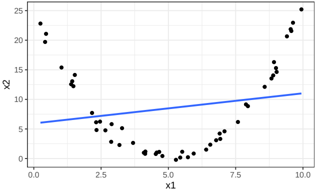

- Interpretable models may also be difficult to interpret, e.g.:
  - Linear model with hundreds of features and interactions
  - Decision trees with huge tree depth

- Often not able to automatically model complex relationships due to limited model flexibility
  e.g., high-order main or interaction effects need to be specified manually in a LM

DISADVANTAGES

- Often certain assumptions about data / model structure required
  $\rightsquigarrow$ If assumptions are wrong, models may perform bad

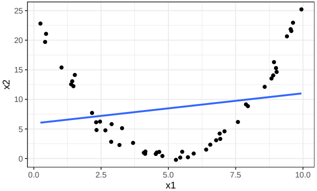

- Interpretable models may also be difficult to interpret, e.g.:
  - Linear model with hundreds of features and interactions
  - Decision trees with huge tree depth

- Often not able to automatically model complex relationships due to limited model flexibility
  e.g., high-order main or interaction effects need to be specified manually in a LM

- Inherently interpretable models do not provide all types of explanations
  $\rightsquigarrow$ Methods providing other types of explanations still useful (e.g., counterfactual explanations)

FURTHER COMMENTS

- Some argue that one should always use interpretable models in the first place ▸ Rudin 2019
  - ...and not try to explain uninterpretable models post-hoc
  - Can sometimes work out by spending enough time and energy on data pre-processing

FURTHER COMMENTS

- Some argue that one should always use interpretable models in the first place ▸ Rudin 2019
  - ...and not try to explain uninterpretable models post-hoc
  - Can sometimes work out by spending enough time and energy on data pre-processing

$\rightsquigarrow$ Drawback: Hard to achieve for data for which end-to-end learning is crucial
(e.g., hard to extract good features for image / text data $\rightsquigarrow$ information loss = bad performance)

FURTHER COMMENTS

- Some argue that one should always use interpretable models in the first place ▸ Rudin 2019
  - . . . and not try to explain uninterpretable models post-hoc
  - Can sometimes work out by spending enough time and energy on data pre-processing

$\rightsquigarrow$ Drawback: Hard to achieve for data for which end-to-end learning is crucial
(e.g., hard to extract good features for image / text data $\rightsquigarrow$ information loss = bad performance)

- Often there is a trade-off between interpretability and model performance

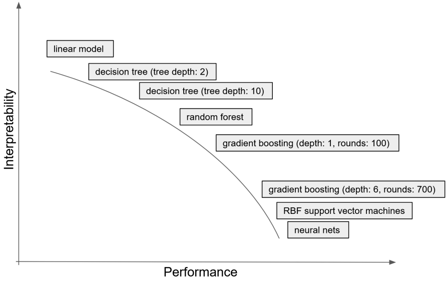

RECOMMENDATION

- Start with most simple model that makes sense for application at hand
- Gradually increase complexity if performance is insufficient
  $\rightsquigarrow$ will usually lower interpretability and require additional interpretation methods
- Choose the most simple, sufficient model (Occam's razor)

Bike Data, 4-fold CV

Model	RMSE	R^2
LM	800.15	0.83
Tree	981.83	0.74
Random Forest	653.25	0.88
Boosting (tuned)	638.42	0.89