


# Interpretable Machine Learning

## Shapley Values



### Learning goals

- Learn what game theory is
- Understand the concept behind cooperative games
- Understand the Shapley value in game theory

# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value

# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players  $P = \{1, \dots, p\}$ , each subset of players  $S \subseteq P$  forms a coalition – each coalition  $S$  achieves a certain payout

# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players  $P = \{1, \dots, p\}$ , each subset of players  $S \subseteq P$  forms a coalition – each coalition  $S$  achieves a certain payout
- A value function  $v : 2^{|P|} \mapsto \mathbb{R}$  maps all  $2^{|P|}$  possible coalitions to their payout (or gain)

# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players  $P = \{1, \dots, p\}$ , each subset of players  $S \subseteq P$  forms a coalition – each coalition  $S$  achieves a certain payout
- A value function  $v : 2^{|P|} \mapsto \mathbb{R}$  maps all  $2^{|P|}$  possible coalitions to their payout (or gain)
- $v(S)$  is the payout of coalition  $S \subseteq P$  (payout of empty coalition must be zero:  $v(\emptyset) = 0$ )

# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players  $P = \{1, \dots, p\}$ , each subset of players  $S \subseteq P$  forms a coalition – each coalition  $S$  achieves a certain payout
- A value function  $v : 2^{|P|} \mapsto \mathbb{R}$  maps all  $2^{|P|}$  possible coalitions to their payout (or gain)
- $v(S)$  is the payout of coalition  $S \subseteq P$  (payout of empty coalition must be zero:  $v(\emptyset) = 0$ )
- As some players contribute more than others, we want to fairly divide the total achievable payout  $v(P)$  among the players according to a player's individual contribution


# COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players  $P = \{1, \dots, p\}$ , each subset of players  $S \subseteq P$  forms a coalition – each coalition  $S$  achieves a certain payout
- A value function  $v : 2^{|P|} \mapsto \mathbb{R}$  maps all  $2^{|P|}$  possible coalitions to their payout (or gain)
- $v(S)$  is the payout of coalition  $S \subseteq P$  (payout of empty coalition must be zero:  $v(\emptyset) = 0$ )
- As some players contribute more than others, we want to fairly divide the total achievable payout  $v(P)$  among the players according to a player's individual contribution
- We call the individual payout per player  $\phi_j, j \in P$  (later: Shapley value)


# COOPERATIVE GAMES WITHOUT INTERACTIONS

Players do not interact  
(payouts ☆ add up in each coalition)





What is a fair payout distribution?

Players do not interact







# COOPERATIVE GAMES WITHOUT INTERACTIONS



⇒ Fair Payouts are Trivial Without Interactions

# COOPERATIVE GAMES WITH INTERACTIONS




⇒ Unclear how to fairly distribute payouts when players interact

# COOPERATIVE GAMES WITH INTERACTIONS

**Question:** What is a fair payout for player “yellow”?

**Idea:** Compute marginal contribution of the player of interest across different coalitions




- Compute the total payout of each coalition
- Compute difference in payouts for each coalition with and without player “yellow” (= marginal contribution)
- Average marginal contributions using appropriate weights


# COOPERATIVE GAMES WITH INTERACTIONS

**Question:** What is a fair payout for player “yellow”?

**Idea:** Compute marginal contribution of the player of interest across different coalitions



**Note:** Each marginal contribution is weighted w.r.t. number of possible orders of its coalition  
 $\leadsto$  More players in  $S \Rightarrow$  more orderings of  $S$



- Compute the total payout of each coalition
- Compute difference in payouts for each coalition with and without player “yellow” (= marginal contribution)
- Average marginal contributions using appropriate weights

# SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let  $v(S \cup \{j\}) - v(S)$  be the marginal contribution of player  $j$  to coalition  $S$   
     $\rightsquigarrow$  measures how much a player  $j$  increases the value of a coalition  $S$

# SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let  $v(S \cup \{j\}) - v(S)$  be the marginal contribution of player  $j$  to coalition  $S$   
     $\rightsquigarrow$  measures how much a player  $j$  increases the value of a coalition  $S$
- Average marginal contributions for all possible coalitions  $S \subseteq P \setminus \{j\}$   
     $\rightsquigarrow$  order of how players join the coalition matters  $\Rightarrow$  different weights depending on size of  $S$

# SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let  $v(S \cup \{j\}) - v(S)$  be the marginal contribution of player  $j$  to coalition  $S$   
     $\rightsquigarrow$  measures how much a player  $j$  increases the value of a coalition  $S$
- Average marginal contributions for all possible coalitions  $S \subseteq P \setminus \{j\}$   
     $\rightsquigarrow$  order of how players join the coalition matters  $\Rightarrow$  different weights depending on size of  $S$
- Shapley value via **set definition** (weighting via multinomial coefficient):

$$\phi_j = \sum_{S \subseteq P \setminus \{j\}} \frac{|S|!(|P| - |S| - 1)!}{|P|!} (v(S \cup \{j\}) - v(S))$$

# SHAPLEY VALUE - ORDER DEFINITION

The Shapley value was introduced as summation over sets  $S \subseteq P \setminus \{j\}$ , but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- $\Pi$ : All possible orders of players (we have  $|P|!$  in total)



# SHAPLEY VALUE - ORDER DEFINITION

The Shapley value was introduced as summation over sets  $S \subseteq P \setminus \{j\}$ , but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- $\Pi$ : All possible orders of players (we have  $|P|!$  in total)
- $S_j^\tau$ : Set of players before player  $j$  in order  $\tau = (\tau^{(1)}, \dots, \tau^{(p)})$  where  $\tau^{(i)}$  is  $i$ -th element
  - $\Rightarrow$  Example: Players 1, 2, 3  $\Rightarrow \Pi = \{(1, 2, 3), (1, 3, 2), (2, 1, 3), (2, 3, 1), (3, 1, 2), (3, 2, 1)\}$ 
    - $\rightsquigarrow$  For order  $\tau = (2, 1, 3)$  and player of interest  $j = 3 \Rightarrow S_j^\tau = \{2, 1\}$
    - $\rightsquigarrow$  For order  $\tau = (3, 1, 2)$  and player of interest  $j = 1 \Rightarrow S_j^\tau = \{3\}$
    - $\rightsquigarrow$  For order  $\tau = (3, 1, 2)$  and player of interest  $j = 3 \Rightarrow S_j^\tau = \emptyset$

# SHAPLEY VALUE - ORDER DEFINITION


The Shapley value was introduced as summation over sets  $S \subseteq P \setminus \{j\}$ , but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- $\Pi$ : All possible orders of players (we have  $|P|!$  in total)
- $S_j^\tau$ : Set of players before player  $j$  in order  $\tau = (\tau^{(1)}, \dots, \tau^{(p)})$  where  $\tau^{(i)}$  is  $i$ -th element
  - $\Rightarrow$  Example: Players 1, 2, 3  $\Rightarrow \Pi = \{(\textcolor{red}{1}, \textcolor{red}{2}, 3), (1, 3, 2), (\textcolor{red}{2}, \textcolor{red}{1}, 3), (2, 3, 1), (3, 1, 2), (3, 2, 1)\}$ 
    - $\rightsquigarrow$  For order  $\tau = (2, 1, 3)$  and player of interest  $j = 3 \Rightarrow S_j^\tau = \{2, 1\}$
    - $\rightsquigarrow$  For order  $\tau = (3, 1, 2)$  and player of interest  $j = 1 \Rightarrow S_j^\tau = \{3\}$
    - $\rightsquigarrow$  For order  $\tau = (3, 1, 2)$  and player of interest  $j = 3 \Rightarrow S_j^\tau = \emptyset$
- Order definition: Marginal contribution of orders that yield set  $S = \{\textcolor{red}{1}, \textcolor{red}{2}\}$  is summed twice
  - $\rightsquigarrow$  In set definition, it has the weight  $\frac{2!(3-2-1)!}{3!} = \frac{2 \cdot 0!}{6} = \frac{2}{6}$


# SHAPLEY VALUE - COMMENTS ON ORDER DEFINITION

- Order and set definition are equivalent
- Reason: The number of orders which yield the same coalition  $S$  is  $|S|!(|P| - |S| - 1)!$ 
  - $\Rightarrow$  There are  $|S|!$  possible orders of players within coalition  $S$
  - $\Rightarrow$  There are  $(|P| - |S| - 1)!$  possible orders of players without  $S$  and  $j$



# SHAPLEY VALUE - COMMENTS ON ORDER DEFINITION

- Order and set definition are equivalent
- Reason: The number of orders which yield the same coalition  $S$  is  $|S|!(|P| - |S| - 1)!$ 
  - $\Rightarrow$  There are  $|S|!$  possible orders of players within coalition  $S$
  - $\Rightarrow$  There are  $(|P| - |S| - 1)!$  possible orders of players without  $S$  and  $j$




- Relevance of the order definition: Approximate Shapley values by sampling permutations  
 $\rightsquigarrow$  randomly sample a fixed number of  $M$  permutations and average them:

$$\phi_j = \frac{1}{M} \sum_{\tau \in \Pi_M} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

where  $\Pi_M \subset \Pi$  is a random subset of  $\Pi$  containing only  $M$  orders of players


# WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION

$|P|! = 6$  orders



# WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION


$|P|! = 6$  orders




via orders

(players before "yellow" player)

$|S| = 2$   
weight =  $1/6$




$|S| = 2$   
weight =  $1/6$



$|S| = 1$   
weight =  $1/6$



$|S| = 1$   
weight =  $1/6$



$|S| = 0$   
weight =  $1/6$




$|S| = 0$   
weight =  $1/6$



# WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION


$|P|! = 6$  orders




via orders

(players before "yellow" player)

$|S| = 2$   
weight =  $1/6$




$|S| = 2$   
weight =  $1/6$



$|S| = 1$   
weight =  $1/6$



$|S| = 1$   
weight =  $1/6$



$|S| = 0$   
weight =  $1/6$




$|S| = 0$   
weight =  $1/6$



via sets


$|S| = 2$   
weight =  $2/6$



$|S| = 1$   
weight =  $1/6$



$|S| = 1$   
weight =  $1/6$




$|S| = 0$   
weight =  $2/6$



# SHAPLEY VALUES - ILLUSTRATION


- Shapley value of player  $j$  is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions






# SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player  $j$  is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 1 enters the coalition



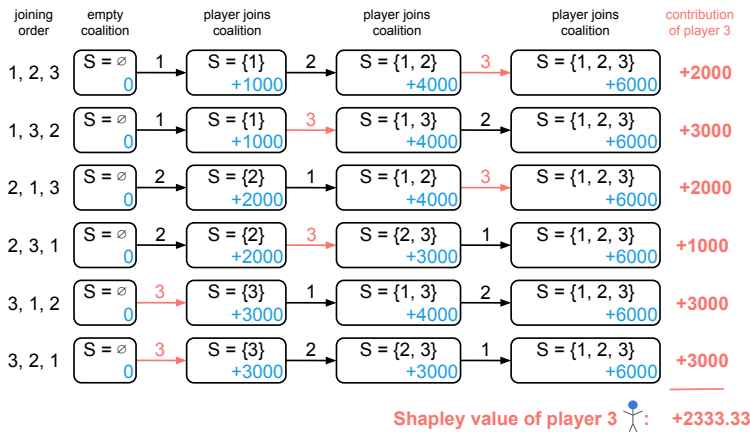
# SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player  $j$  is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 2 enters the coalition




# SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player  $j$  is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 3 enters the coalition



# SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player  $j$  is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions



# AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function  $v$ :

- **Efficiency:** Player contributions add up to the total payout of the game:  $\sum_{j=1}^p \phi_j = v(P)$

# AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function  $v$ :

- **Efficiency:** Player contributions add up to the total payout of the game:  $\sum_{j=1}^P \phi_j = v(P)$
- **Symmetry:** Players  $j, k \in P$  who contribute the same to any coalition get the same payout:  
If  $v(S \cup \{j\}) = v(S \cup \{k\})$  for all  $S \subseteq P \setminus \{j, k\}$ , then  $\phi_j = \phi_k$

# AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function  $v$ :

- **Efficiency:** Player contributions add up to the total payout of the game:  $\sum_{j=1}^P \phi_j = v(P)$
- **Symmetry:** Players  $j, k \in P$  who contribute the same to any coalition get the same payout:  
If  $v(S \cup \{j\}) = v(S \cup \{k\})$  for all  $S \subseteq P \setminus \{j, k\}$ , then  $\phi_j = \phi_k$
- **Dummy/Null Player:** Payout is 0 for players who don't contribute to the value of any coalition:  
If  $v(S \cup \{j\}) = v(S) \quad \forall \quad S \subseteq P \setminus \{j\}$ , then  $\phi_j = 0$

# AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function  $v$ :

- **Efficiency:** Player contributions add up to the total payout of the game:  $\sum_{j=1}^P \phi_j = v(P)$
- **Symmetry:** Players  $j, k \in P$  who contribute the same to any coalition get the same payout:  
If  $v(S \cup \{j\}) = v(S \cup \{k\})$  for all  $S \subseteq P \setminus \{j, k\}$ , then  $\phi_j = \phi_k$
- **Dummy/Null Player:** Payout is 0 for players who don't contribute to the value of any coalition:  
If  $v(S \cup \{j\}) = v(S) \quad \forall \quad S \subseteq P \setminus \{j\}$ , then  $\phi_j = 0$
- **Additivity:** For a game  $v$  with combined payouts  $v(S) = v_1(S) + v_2(S)$ , the payout is the sum of payouts:  $\phi_{j,v} = \phi_{j,v_1} + \phi_{j,v_2}$