



Interpretable Machine Learning

LIME




Learning goals

- Understand motivation for LIME
- Develop a mathematical intuition




LIME

- Local Interpretable Model-agnostic Explanations (LIME) assume that even if a ML model is very complex, the local prediction can be described with a simpler model




LIME

- Local Interpretable Model-agnostic Explanations (LIME) assume that even if a ML model is very complex, the local prediction can be described with a simpler model
- LIME explains **individual** predictions of **any** black-box model by approximating the model **locally** with an interpretable model




LIME

- Local Interpretable Model-agnostic Explanations (LIME) assume that even if a ML model is very complex, the local prediction can be described with a simpler model
- LIME explains **individual** predictions of **any** black-box model by approximating the model **locally** with an interpretable model
- Called local surrogate models \rightsquigarrow often inherently interpretable models such as linear models or classification/regression trees are chosen




LIME

- Local Interpretable Model-agnostic Explanations (LIME) assume that even if a ML model is very complex, the local prediction can be described with a simpler model
- LIME explains **individual** predictions of **any** black-box model by approximating the model **locally** with an interpretable model
- Called local surrogate models \rightsquigarrow often inherently interpretable models such as linear models or classification/regression trees are chosen
- LIME should answer why a ML model predicted \hat{y} for input \mathbf{x}



LIME

- Local Interpretable Model-agnostic Explanations (LIME) assume that even if a ML model is very complex, the local prediction can be described with a simpler model
- LIME explains **individual** predictions of **any** black-box model by approximating the model **locally** with an interpretable model
- Called local surrogate models \rightsquigarrow often inherently interpretable models such as linear models or classification/regression trees are chosen
- LIME should answer why a ML model predicted \hat{y} for input \mathbf{x}
- LIME is model-agnostic and can handle tabular, image and text data



LIME: CHARACTERISTICS


Definition:

LIME provides a local explanation for a black-box model \hat{f} in form of a model $\hat{g} \in \mathcal{G}$ with \mathcal{G} as the class of potential (interpretable) models

Model g should have two characteristics:

- 1 **Interpretable**: relation between the input variables and the response are easy to understand
- 2 **Locally faithful / Fidelity**: similar behavior as \hat{f} in the vicinity of the obs. being predicted

Formally, we want to receive a model \hat{g} with **minimal complexity and maximal local-fidelity**




MODEL COMPLEXITY

We can measure the complexity of a model \hat{g} using a complexity measure $J(\hat{g})$

Example: Linear model

- Let $\mathcal{G} = \{g : \mathcal{X} \rightarrow \mathbb{R} \mid g(\mathbf{x}) = s(\boldsymbol{\theta}^\top \mathbf{x})\}$ be the class of linear models
- $s(\cdot)$: identity function for linear regression or logistic sigmoid function for logistic regression

$\rightsquigarrow J(g) = \sum_{j=1}^p \mathcal{I}_{\{\theta_j \neq 0\}}$ could be the L_0 loss, i.e., the number of non-zero coefficients



MODEL COMPLEXITY


We can measure the complexity of a model \hat{g} using a complexity measure $J(\hat{g})$

Example: Linear model

- Let $\mathcal{G} = \{g : \mathcal{X} \rightarrow \mathbb{R} \mid g(\mathbf{x}) = s(\theta^\top \mathbf{x})\}$ be the class of linear models
 - $s(\cdot)$: identity function for linear regression or logistic sigmoid function for logistic regression
- $\rightsquigarrow J(g) = \sum_{j=1}^p \mathcal{I}_{\{\theta_j \neq 0\}}$ could be the L_0 loss, i.e., the number of non-zero coefficients


Example: Tree

- Let $\mathcal{G} = \left\{ g : \mathcal{X} \rightarrow \mathbb{R} \mid g(\mathbf{x}) = \sum_{m=1}^M c_m \mathcal{I}_{\{\mathbf{x} \in Q_m\}} \right\}$ be the class of trees
i.e., the class of additive models (e.g., constant c_m) over the leaf-rectangles Q_m
- $\rightsquigarrow J(g)$ could measure the number of terminal/leaf nodes



LOCAL MODEL FIDELITY

- g is locally faithful to \hat{f} w.r.t. \mathbf{x} if for $\mathbf{z} \in \mathcal{Z} \subseteq \mathbb{R}^p$ close to \mathbf{x} , predictions of $\hat{g}(\mathbf{z})$ are close to $\hat{f}(\mathbf{z})$
- In an optimization task: the closer \mathbf{z} is to \mathbf{x} , the closer $\hat{g}(\mathbf{z})$ should be to $\hat{f}(\mathbf{z})$




LOCAL MODEL FIDELITY

- g is locally faithful to \hat{f} w.r.t. \mathbf{x} if for $\mathbf{z} \in \mathcal{Z} \subseteq \mathbb{R}^p$ close to \mathbf{x} , predictions of $\hat{g}(\mathbf{z})$ are close to $\hat{f}(\mathbf{z})$
- In an optimization task: the closer \mathbf{z} is to \mathbf{x} , the closer $\hat{g}(\mathbf{z})$ should be to $\hat{f}(\mathbf{z})$
- Two required measures:
 - ❶ A proximity (similarity) measure $\phi_{\mathbf{x}}(\mathbf{z})$ between \mathbf{z} and \mathbf{x} , e.g. the exponential kernel:

$$\phi_{\mathbf{x}}(\mathbf{z}) = \exp(-d(\mathbf{x}, \mathbf{z})^2 / \sigma^2)$$

with σ as the kernel width and d as the Euclidean distance (numeric features) or the Gower distance (mixed features)



LOCAL MODEL FIDELITY


- g is locally faithful to \hat{f} w.r.t. \mathbf{x} if for $\mathbf{z} \in \mathcal{Z} \subseteq \mathbb{R}^p$ close to \mathbf{x} , predictions of $\hat{g}(\mathbf{z})$ are close to $\hat{f}(\mathbf{z})$
- In an optimization task: the closer \mathbf{z} is to \mathbf{x} , the closer $\hat{g}(\mathbf{z})$ should be to $\hat{f}(\mathbf{z})$
- Two required measures:
 - ❶ A proximity (similarity) measure $\phi_{\mathbf{x}}(\mathbf{z})$ between \mathbf{z} and \mathbf{x} , e.g. the exponential kernel:

$$\phi_{\mathbf{x}}(\mathbf{z}) = \exp(-d(\mathbf{x}, \mathbf{z})^2 / \sigma^2)$$

with σ as the kernel width and d as the Euclidean distance (numeric features) or the Gower distance (mixed features)

- ❷ A distance measure or loss function $L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z}))$, e.g. the L_2 loss/squared error

$$L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z})) = (\hat{g}(\mathbf{z}) - \hat{f}(\mathbf{z}))^2$$



LOCAL MODEL FIDELITY

- g is locally faithful to \hat{f} w.r.t. \mathbf{x} if for $\mathbf{z} \in \mathcal{Z} \subseteq \mathbb{R}^p$ close to \mathbf{x} , predictions of $\hat{g}(\mathbf{z})$ are close to $\hat{f}(\mathbf{z})$
- In an optimization task: the closer \mathbf{z} is to \mathbf{x} , the closer $\hat{g}(\mathbf{z})$ should be to $\hat{f}(\mathbf{z})$
- Two required measures:
 - ❶ A proximity (similarity) measure $\phi_{\mathbf{x}}(\mathbf{z})$ between \mathbf{z} and \mathbf{x} , e.g. the exponential kernel:

$$\phi_{\mathbf{x}}(\mathbf{z}) = \exp(-d(\mathbf{x}, \mathbf{z})^2 / \sigma^2)$$


with σ as the kernel width and d as the Euclidean distance (numeric features) or the Gower distance (mixed features)

- ❷ A distance measure or loss function $L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z}))$, e.g. the L_2 loss/squared error

$$L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z})) = (\hat{g}(\mathbf{z}) - \hat{f}(\mathbf{z}))^2$$

- Given points \mathbf{z} , we can measure local fidelity of g with respect to \hat{f} in terms of a weighted loss

$$L(\hat{f}, g, \phi_{\mathbf{x}}) = \sum_{\mathbf{z} \in \mathcal{Z}} \phi_{\mathbf{x}}(\mathbf{z}) L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z}))$$




MINIMIZATION TASK

- Optimization objective of LIME:

$$\arg \min_{g \in \mathcal{G}} L(\hat{f}, \hat{g}, \phi_{\mathbf{x}}) + J(g)$$


- In practice:
 - LIME only optimizes $L(\hat{f}, \hat{g}, \phi_{\mathbf{x}})$ (model-fidelity)
 - Users decide threshold on model complexity $J(g)$ beforehand
- Goal: **model-agnostic** explainer
 - ↪ optimize $L(\hat{f}, \hat{g}, \phi_{\mathbf{x}})$ without making any assumptions about \hat{f}
 - ↪ learn \hat{g} only approximately



LIME ALGORITHM: OUTLINE

Input:

- Pre-trained model \hat{f}
- Observation \mathbf{x} whose prediction $\hat{f}(\mathbf{x})$ we want to explain
- Model class \mathcal{G} for local surrogate (to limit the complexity of the explanation)




LIME ALGORITHM: OUTLINE

Input:

- Pre-trained model \hat{f}
- Observation \mathbf{x} whose prediction $\hat{f}(\mathbf{x})$ we want to explain
- Model class \mathcal{G} for local surrogate (to limit the complexity of the explanation)

Algorithm:



- 1 Independently sample new points $\mathbf{z} \in \mathcal{Z}$
- 2 Retrieve predictions $\hat{f}(\mathbf{z})$ for obtained points \mathbf{z}
- 3 Weight $\mathbf{z} \in \mathcal{Z}$ by their proximity $\phi_{\mathbf{x}}(\mathbf{z})$
- 4 Train an interpretable surrogate model g on weighted data points $\mathbf{z} \in \mathcal{Z}$
 \rightsquigarrow predictions $\hat{f}(\mathbf{z})$ are the target of this model
- 5 Return the interpretable model \hat{g} as the explainer



LIME ALGORITHM: EXAMPLE

Illustration of LIME based on a classification task:

- Light/dark gray background: prediction surface of a classifier
- Yellow point: \mathbf{x} to be explained
- \mathcal{G} : class of logistic regression models



LIME ALGORITHM: EXAMPLE (STEP 1+2: SAMPLING)

► Ribeiro. 2016

Strategies for sampling:

- Uniformly sample new points from the feasible feature range
- Use the training data set with or without perturbations
- Draw samples from the estimated univariate distribution of each feature
- Create an equidistant grid over the supported feature range




Figure: Uniformly sampled





Figure: Equidistant grid





LIME ALGORITHM: EXAMPLE (STEP 3: PROXIMITY)

► Ribeiro. 2016

In this example, we use the exponential kernel defined on the Euclidean distance d

$$\phi_{\mathbf{x}}(\mathbf{z}) = \exp(-d(\mathbf{x}, \mathbf{z})^2 / \sigma^2).$$



LIME ALGORITHM: EXAMPLE (STEP 4: SURROGATE)

► Ribeiro. 2016

In our example, we fit a **logistic regression** model (consequently, $L(\hat{f}(\mathbf{z}), \hat{g}(\mathbf{z}))$ is the Bernoulli loss)

